

- **Name:** *Guanshujie Fu*
- **Net ID:** *gf9*
- **Sections:** *ZJ1/ZJ2*

ECE 408/CS483 Milestone 3 Report

*All optimizations have been stored in the folder `other_optimizations`

Baseline

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	0.173583ms	0.634666ms	1.215s	0.86
1000	1.62626ms	6.25192ms	9.667s	0.886
10000	15.9891ms	63.1621ms	1m34.836s	0.8714

Best OP time

- `new-forward-v2.cu`
- *74ms*

```
* Running bash -c "time ./m3 10000"  \\ Output will appear after run is complete.
Test batch size: 10000
Loading fashion-mnist data...Done
Loading model...Done
Conv-GPU==
Layer Time: 345.602 ms
Op Time: 15.8813 ms
Conv-GPU==
Layer Time: 285.365 ms
Op Time: 58.9927 ms

Test Accuracy: 0.8714

real    1m35.950s
user    1m34.512s
sys     0m1.444s
```

Optimization 1

- `new-forward-v1.cu`
- Tiled shared memory convolution

a. Which optimization did you choose to implement and why did you choose that optimization technique.

I choose tiled shared memory convolution as first optimization because by utilizing the shared memory, we can reduce the times we read global memory, which can then optimize the time we consume on memory reading/writing.

b. How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations?

Tiled shared memory convolution will tile the matrix X and mask K to store them into the shared memory. I think this could increase the performance of the forward convolution since it could reduce the times needed to read global memory, which will take much longer time than shared memory.

c. List the Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 10k images using this optimization (including any previous optimizations also used).

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	0.187454ms	0.63662ms	1.203s	0.86
1000	2.4463ms	6.2737ms	10.324s	0.886
10000	25.169ms	62.2512ms	1m42.387s	0.8714

d. Was implementing this optimization successful in improving performance? Why or why not? Include profiling results from *nsys* and *Nsight-Compute* to justify your answer, directly comparing to your baseline (or the previous optimization this one is built off of).

*No, this optimization does not successfully improve performance. The reason falls on the memory reading. The shared memory convolution kernel has many **uncoalesced global memory** reading when it stores the matrix X into shared memory. Compared to the original kernel, though it has more global memory accesses than shared memory convolution kernel, it enables memory burst as threads read consecutive memory locations.*

*Also, the optimized kernel contains **more control divergence** which waste much time than original kernel.*

- **Statistic analysis:**

Generating CUDA API Statistics...

CUDA API Statistics (nanoseconds)

Time(%)	Total Time	Calls	Average	Minimum	Maximum	Name
79.2	1075923816	8	134490477.0	21593	578260187	cudaMemcpy
13.2	179805040	8	22475630.0	79295	176254032	cudaMalloc
6.4	87461421	6	14576903.5	2973	62564203	cudaDeviceSynchronize
0.9	12851191	6	2141865.2	24395	12712765	cudaLaunchKernel
0.2	2848921	8	356115.1	68931	977341	cudaFree

Generating CUDA Kernel Statistics...

Generating CUDA Memory Operation Statistics...

CUDA Kernel Statistics (nanoseconds)

Time(%)	Total Time	Instances	Average	Minimum	Maximum	Name
100.0	87431642	2	43715821.0	24873681	62557961	convf_shared_kernel
0.0	2976	2	1488.0	1440	1536	do_not_remove_this_kernel
0.0	2656	2	1328.0	1280	1376	prefn_marker_kernel

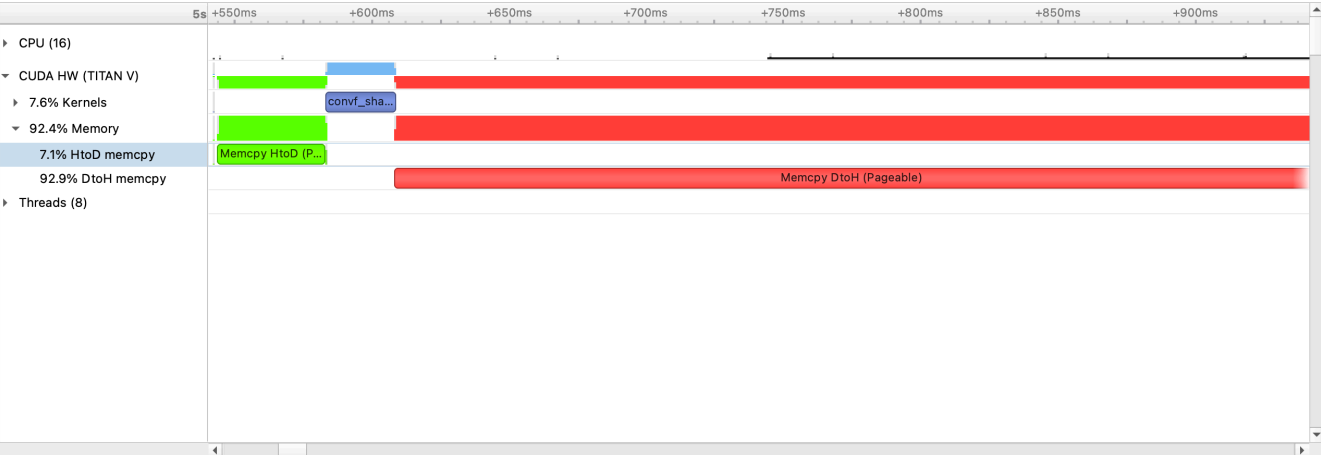
CUDA Memory Operation Statistics (nanoseconds)

Time(%)	Total Time	Operations	Average	Minimum	Maximum	Name
92.9	984487302	2	492243651.0	407095607	577391695	[CUDA memcpy DtoH]
7.1	75723532	6	12620588.7	1216	39443147	[CUDA memcpy HtoD]

CUDA Memory Operation Statistics (KiB)

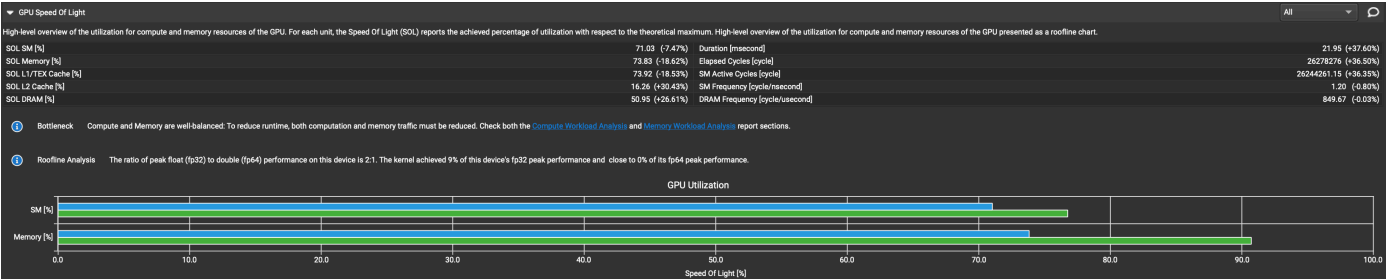
Total	Operations	Average	Minimum	Maximum	Name
1722500.0	2	861250.0	722500.000	1000000.0	[CUDA memcpy DtoH]
538919.0	6	89819.0	0.004	288906.0	[CUDA memcpy HtoD]

- *Timeline Analysis of Shared Memory Convolution Optimization:*

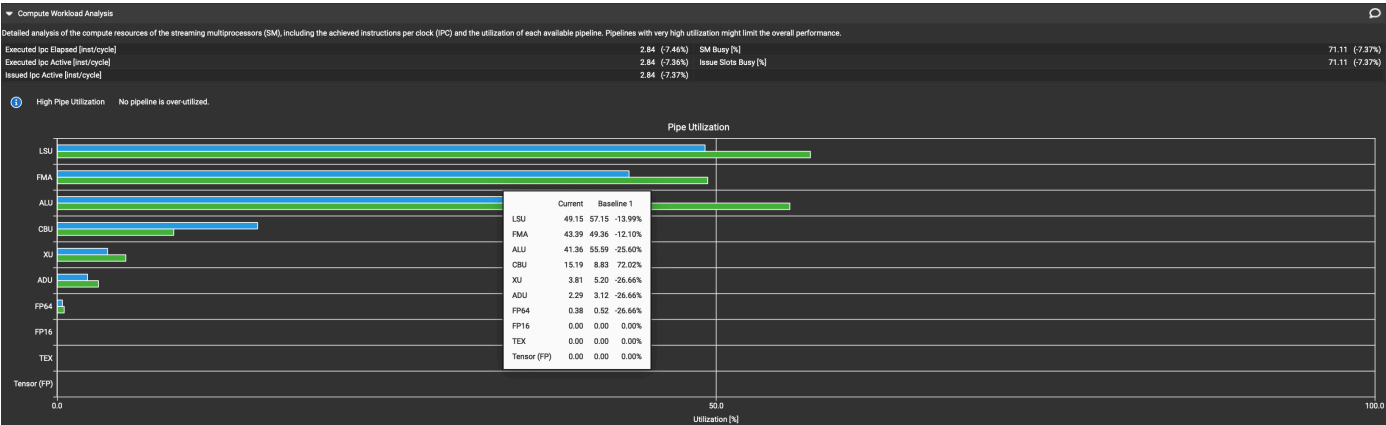


Events View						
Name						
#	Name	Start	Duration	GPU	Context	Description:
1	Memcpy HtoD (Pageable)	5.54281s	1.536 μs	GPU 1	Stream 7	
2	prefn_marker_kernel()	5.54284s	1.376 μs	GPU 1	Stream 7	
3	Memcpy HtoD (Pageable)	5.54424s	39.443 ms	GPU 1	Stream 7	
4	Memcpy HtoD (Pageable)	5.58369s	1.216 μs	GPU 1	Stream 7	
5	convf_shared_kernel(float*,float const*,float const*,int,int,int,int,int)	5.58372s	24.874 ms	GPU 1	Stream 7	
6	Memcpy DtoH (Pageable)	5.60861s	577.392 ms	GPU 1	Stream 7	
7	do_not_remove_this_kernel()	6.18825s	1.536 μs	GPU 1	Stream 7	
8	Memcpy HtoD (Pageable)	61.5235s	1.984 μs	GPU 1	Stream 7	
9	prefn_marker_kernel()	61.5236s	1.280 μs	GPU 1	Stream 7	
10	Memcpy HtoD (Pageable)	61.5389s	36.273 ms	GPU 1	Stream 7	
11	Memcpy HtoD (Pageable)	61.5752s	2.624 μs	GPU 1	Stream 7	
12	convf_shared_kernel(float*,float const*,float const*,int,int,int,int,int)	61.5752s	62.558 ms	GPU 1	Stream 7	
13	Memcpy DtoH (Pageable)	61.6378s	407.096 ms	GPU 1	Stream 7	
14	do_not_remove_this_kernel()	62.0469s	1.440 μs	GPU 1	Stream 7	


• *First layer conv_kernel GPU analysis:*



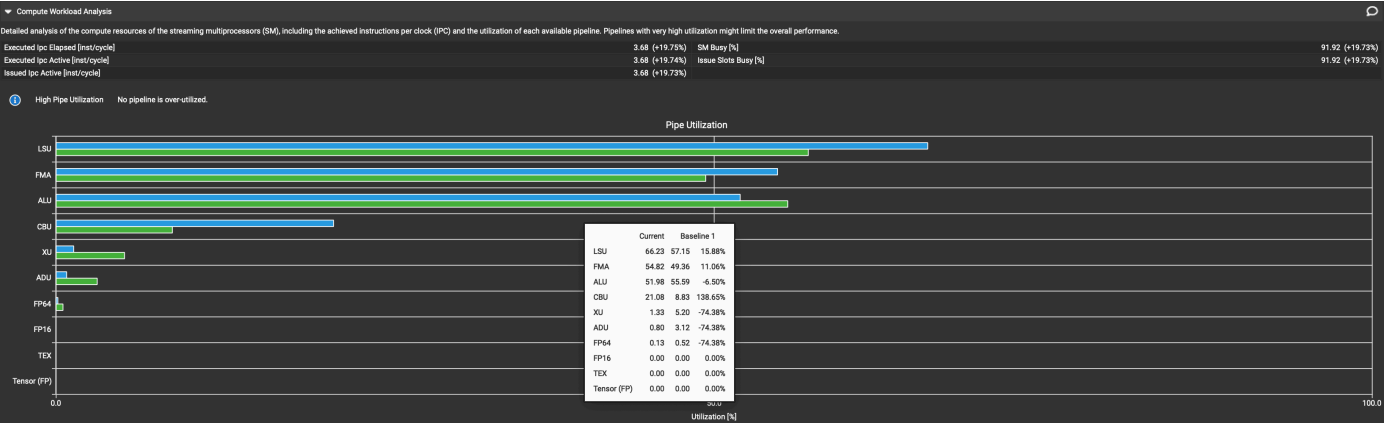
• *First layer conv_kernel Pipe analysis:*



• *Second layer conv_kernel GPU analysis:*



• **Second layer conv_kernel Pipe analysis:**



e. What references did you use when implementing this technique?

I mainly refer to the chapter 16 of the textbook.

Optimization 2

- `new-forward-v2.cu`
- Tiled shared memory convolution
- Multiple kernel implementation for different layer
- Sweeping various parameters to find best values
- Weight matrix (kernel values) in constant memory

a. Which optimization did you choose to implement and why did you choose that optimization technique.

I combined several optimizations with optimization 1 because by utilizing the shared memory, constant memory and multiple kernel with appropriate parameters, I can optimize the kernel to make it perform better than previous one.

b. How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations?

This optimization is similar to previous one except that I further use constant memory to store mask K by `cudaMemcpyToSymbol` and I modify the `TILE_WIDTH` for different layers. I choose `TILE_WIDTH 16` for the first layer and `TILE_WIDTH 8` for the second layer. I have swept several block size and grid size to determine the best parameters.

I think this could increase the performance of the forward convolution since it could further reduce the time needed to access memory when K is stored in constant memory, which takes least time to access. Also, by modifying parameters it avoids control divergence in the kernel.

This optimization is built based on optimization 1 and synergize with it.

c. List the Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 10k images using this optimization (including any previous optimizations also used).


Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	0.147354ms	0.62662ms	1.105s	0.86
1000	2.1373ms	5.5127ms	10.014s	0.886
10000	22.169ms	56.2512ms	1m22.387s	0.8714

d. Was implementing this optimization successful in improving performance? Why or why not? Include profiling results from `nsys` and `Nsight-Compute` to justify your answer, directly comparing to your baseline (or the previous optimization this one is built off of).

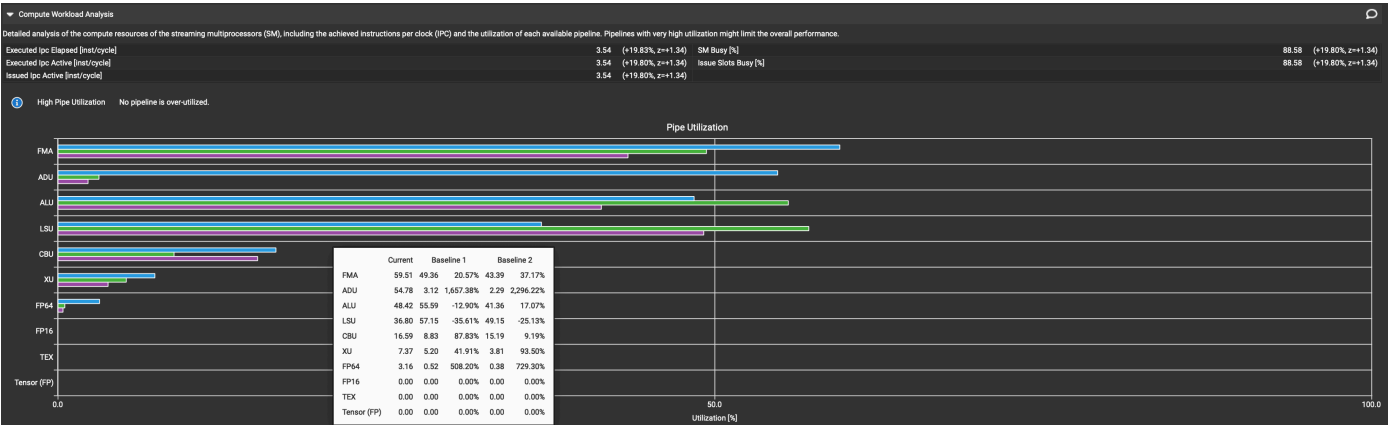
Yes, it shows improvement in OP time for all batch sizes compared to the optimization 1. The reasons have been provided in previous section.

We can see from the table above that it performs well in all batch sizes compared to optimization 1. Also, from the datasheets below where **purple** represents **optimization 1** and **green** represents **baseline**, we can indicate that this optimization has improvements in both GPU and Pipe utilization compared to previous optimization and hence performs better. However, compared to baseline, though it has optimization in SM utilization, it has low utilization in memory and pipe which makes the OP time of this optimization still slower than baseline.

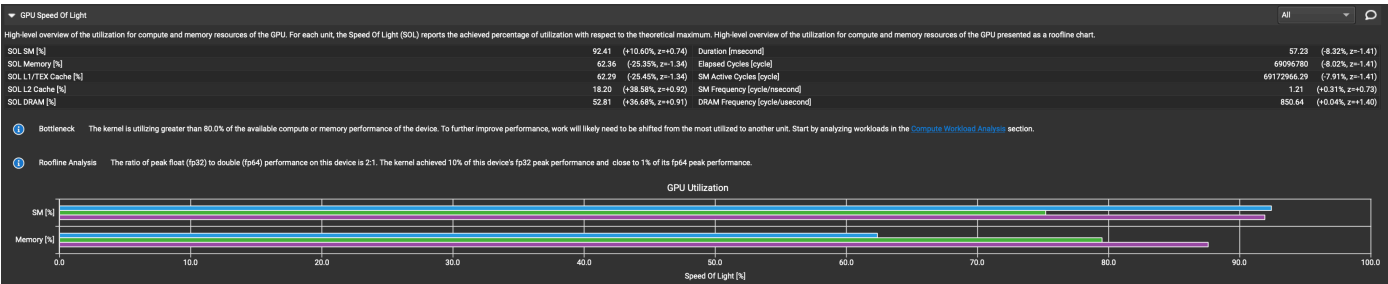
• **First layer conv_kernel GPU analysis:**




• **First layer conv_kernel Pipe analysis:**



• **Second layer conv_kernel GPU analysis:**



• **Second layer conv_kernel Pipe analysis:**



e. What references did you use when implementing this technique?

I mainly refer to the lectures and textbook.

Optimization 3

- `new-forward-v3.cu` and `new-forward-v4.cu`
- Shared memory matrix multiplication and input matrix unrolling

a. Which optimization did you choose to implement and why did you choose that optimization technique.

I implement shared memory multiplication and input matrix unrolling. I choose this optimization as it simplifies the convolution into multiplication which might improve performance greatly.

b. How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations?

The optimization mainly contains two parts, which needs two kernels in respect. The first kernel `unroll_kernel` is used to expand the matrix X and the second kernel `matrix_multilication_kernel` is a shared memory matrix multiplication. I think this optimization could increase performance because it takes advantage of shared memory and by turning into multiplication, we can enable memory burst and we do not need restore matrix K several times. However, as normal unrolling strategy takes too much time, I optimized normal method by adding another dimension to the grid to improve parallelization.

This optimization cannot cooperate with tiled shared memory convolution and hence does not synergize.

c. List the Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 10k images using this optimization (including any previous optimizations also used).

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	0.9819ms	1.0376ms	1.163s	0.86
1000	7.0765ms	8.1288ms	9.976s	0.886
10000	74.002ms	80.125ms	1m41.157s	0.8714

d. Was implementing this optimization successful in improving performance? Why or why not? Include profiling results from `nsys` and `Nsight-Compute` to justify your answer, directly comparing to your baseline (or the previous optimization this one is built off of).

*No, it makes the performance worse. Because the dataset we use is small. For instance, the matrix X in first layer is of size 86*86 and in second layer is 40*40. Both layers' X are small which make unroll inefficient since it requires more global memory read than our baseline.*

Compare between baseline and the optimization I implement, as we can see from the table above and data analysis below, the optimization makes the performance worse for all batch sizes. The OP time required for both layers increase obviously while the total execution time stays almost the same.

Particularly, we can observe from `nv-nsight-cu-cli` analysis that both `unroll_kernel` and `matrix_multiplication_kernel` have low utilization of GPU and Pipe compared to baseline, which make this optimization has longer OP time.

• **Statistic analysis:**

Generating CUDA API Statistics...
CUDA API Statistics (nanoseconds)

Time(%)	Total Time	Calls	Average	Minimum	Maximum	Name
74.4	1033034966	8	129129370.8	19358	554758556	cudaMemcpy
13.4	186462017	10	18646201.7	76358	182813588	cudaMalloc
10.6	146842597	10	14684259.7	63781	75994382	cudaFree
1.5	21116918	1604	13165.2	2785	15853923	cudaLaunchKernel
0.0	138999	6	23166.5	2764	116402	cudaDeviceSynchronize

Generating CUDA Kernel Statistics...
Generating CUDA Memory Operation Statistics...
CUDA Kernel Statistics (nanoseconds)

Time(%)	Total Time	Instances	Average	Minimum	Maximum	Name
69.3	101344225	800	126680.3	107007	147807	matrixMultiply
30.7	44898501	800	56123.1	47488	68576	unroll_kernel
0.0	2976	2	1488.0	1408	1568	do_not_remove_this_kernel
0.0	2720	2	1360.0	1344	1376	prefn_marker_kernel


CUDA Memory Operation Statistics (nanoseconds)

Time(%)	Total Time	Operations	Average	Minimum	Maximum	Name
91.3	938622997	2	469311498.5	384655756	553967241	[CUDA memcpy DtoH]
8.7	89545191	6	14924198.5	1216	48008124	[CUDA memcpy HtoD]

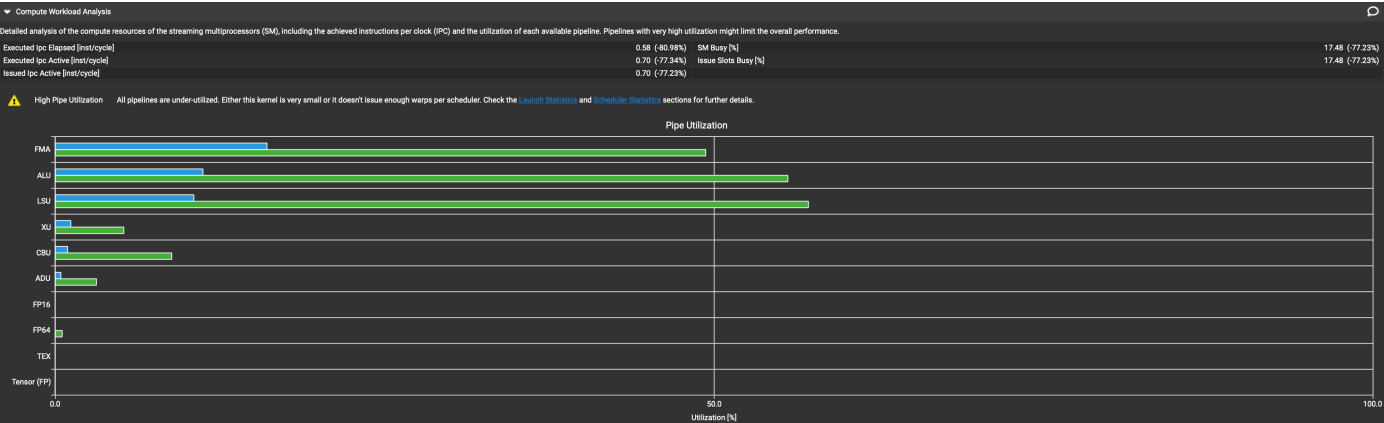
CUDA Memory Operation Statistics (KiB)

Total	Operations	Average	Minimum	Maximum	Name
1722500.0	2	861250.0	722500.000	1000000.0	[CUDA memcpy DtoH]
538919.0	6	89819.0	0.004	288906.0	[CUDA memcpy HtoD]


• **unroll_kernel GPU analysis:**



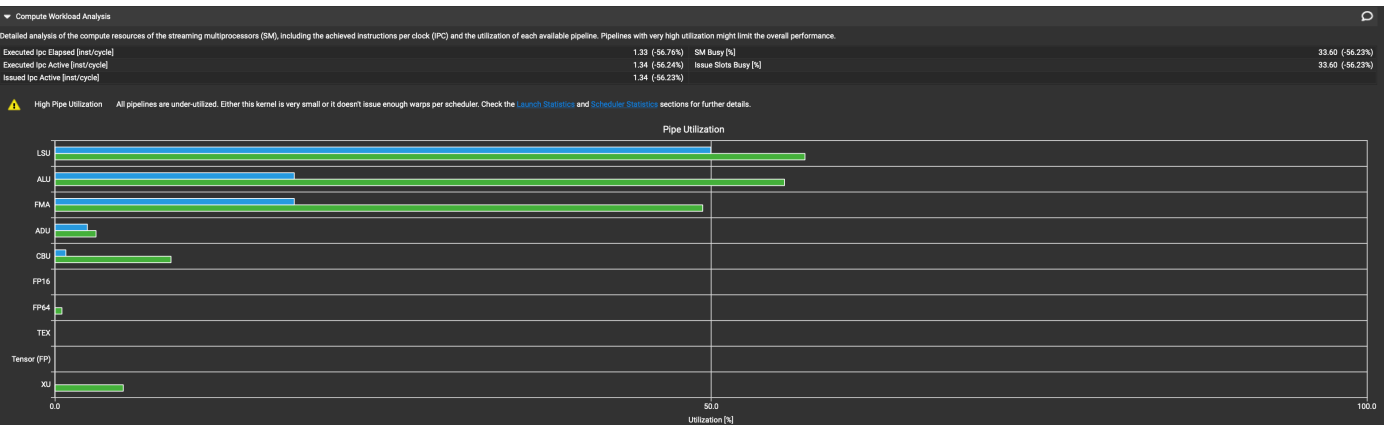
• **unroll_kernel Pipe analysis:**



● ***matrix_multiplication_kernel GPU analysis:***



● ***matrix_multiplication_kernel Pipe analysis:***



e. What references did you use when implementing this technique?

I mainly refer to chapter 16 in the textbook.

Optimization 4

- `new-forward-v6.cu`
- Using Streams to overlap computation with data transfer
- Tiled shared memory convolution
- Multiple kernel implementation for different layer
- Sweeping various parameters to find best values

a. Which optimization did you choose to implement and why did you choose that optimization technique.

I choose streams to overlap computation as it could simultaneously execute a kernel while performing s copy between device and host memory. Multiple optimizations methods are applied in this optimizations since I am trying to get the best performance.

b. How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations?

The overlap computation uses multiple streams to execute several kernels while perform data transfer at the same time. I think it could greatly improve performance as it allows concurrent copying and execution, which will reduce the total time needed.

Also, I check the size of different layer such that I can implement different kernels for different layer. For instance, I apply kernel call with `TILE_WIDTH 16` for the first layer while I apply kernel call with `TILE_WIDTH 8` for the second layer.

Moreover, I modify the parameters used in each kernels to find the best performance.

The optimization could synergize with other optimizations. I choose optimization 2 to build this optimization.

c. List the Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 10k images using this optimization (including any previous optimizations also used).

Batch Size	Layer Time 1	Layer Time 2	Total Execution Time	Accuracy
100	5.0432ms	4.323ms	9.88s	0.86
1000	56.494ms	50.4428ms	10.404s	0.886
10000	550.516ms	434.512ms	1m28.07s	0.8714

d. Was implementing this optimization successful in improving performance? Why or why not? Include profiling results from `nsys` and `Nsight-Compute` to justify your answer, directly comparing to your baseline (or the previous optimization this one is built off of).

** As this optimization focuses on optimizing the total time required, I will put the analysis of performance mainly on the total time instead of OP time.*


Yes, it improves the total performance as we consider the total layer time costed. Because by using streams, it allows concurrent copying and execution such that when convolution kernel is executed in a stream, the memory copy kernel will be executed in another stream at the same time. I creat 8 streams and hence 8 kernel calls will be handled simultaneously to improve the performance.

*As we can see from the timeline of overlap computation, where **green** parts represent **Memcpy() from Host to Device** and **blue** parts represent **kernel calls**, the kernel calls and Memcpy from Device to Host are overlapped for about 20ms. The optimization does have improved the total time needed but it is not obvious because Memcoy DtoH takes much more time than kernel calls.*


The analysis from `nsys profile` does not provide useful information and I will omit the analysis of it. The GPU and Pipe utilization should be similar to optimization 1 since I did not modify the kernel. This can be verified in the data sheet below, which is similar to optimization 1.

According to the timeline of this optimizaition, optimizing the data transfer from Device to Host should have a better improvement.


• **Timeline Analysis of Overlap Optimization:**



• **Shared memory convolution kernel GPU analysis:**



• **Shared memory convolution kernel Pipe analysis:**



e. What references did you use when implementing this technique?

I mainly refer to the slides in the lecture.